

THE C10K problem

concurrently handling ten thousand connections



The Clok problem is the problem of optimising network sockets to handle a large number of clients at the same time.



The C10K problem is the problem of optimising network ~~sockets~~ to handle a large number of clients at the same time.



The Clok problem is the problem of optimising network sockets / network bandwidth / RAM / CPU / OS to handle a large number of clients at the same time.

Little's Law

$$L = \lambda W$$



The long-term average number of customers in a **stable** system L is equal to the long-term average effective arrival rate λ , multiplied by the average time a customer spends in the system W expressed algebraically: $L = \lambda W$.

But what does it mean for us

$$L = \lambda W \rightarrow \lambda = L/W$$

L - system capacity

λ - average request arrival rate

W - average request processing time

To handle more requests, we need to increase L, our capacity, or decrease W, our processing time, or latency.

What Dominates the Capacity (L)

- number of concurrent TCP connections
- network bandwidth
- RAM
- CPU
- OS

BUT does it?

number of concurrent TCP connections

- Server can support several tens-of-thousands of concurrent TCP connections
- Some have had success maintaining over 2 million open connections:
<http://blog.whatsapp.com/196/1-million-is-so-2011>

network bandwidth

- Requests travelling back and forth are no more than a few kilobytes in length (well under 1MB).
- Given today's high-bandwidth LANs, our network could support anywhere between 100K to over a million concurrent requests.


RAM

- how much memory each request consumes
 - let's assume we can keep it below 1MB
- with 500GB of RAM it makes more than 0.5 million **simultaneous** requests
- with 32GB of RAM which is normal even for laptops it makes hundred thousands **concurrent** requests

CPU


- It depends on the application logic ;)
- Most of the time our requests just wait for:
 - microservices
 - database
 - files
- Usually this number is anywhere between several hundreds of thousands and several millions...

in practice productions systems rarely report
CPU as their bottleneck



L = 100K - 1M ???
...

L = 100K - 1M ???




OS

- Number of concurrent request we can handle is also limited by the number of threads the OS can handle
- This number is somewhere between 500 and 15K
- OK, I know there is NPTL for Linux but let's leave it for now...

<http://paultyma.blogspot.in/2008/03/writing-java-multithreaded-servers.html>


$L = 10K \quad ???$
...

$L = 10K \dots$



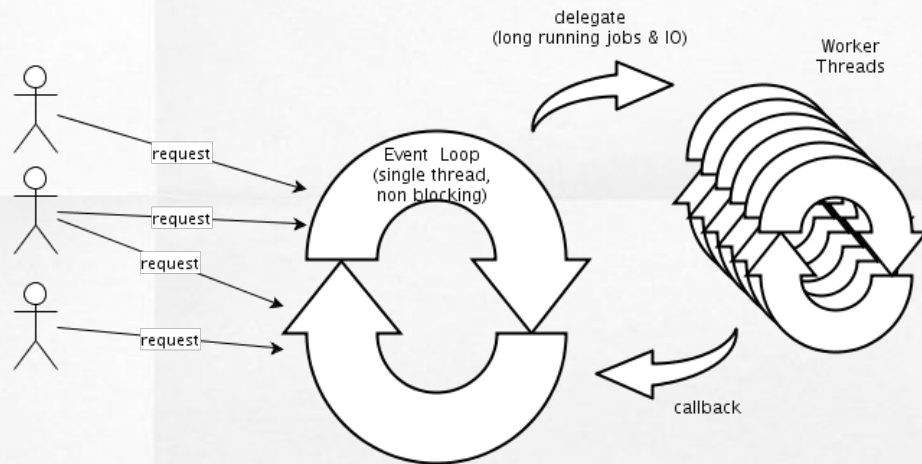
Enterprise architecture ;)

- clusters
- load balancers
- caches
- session replication
- sticky sessions
- ...



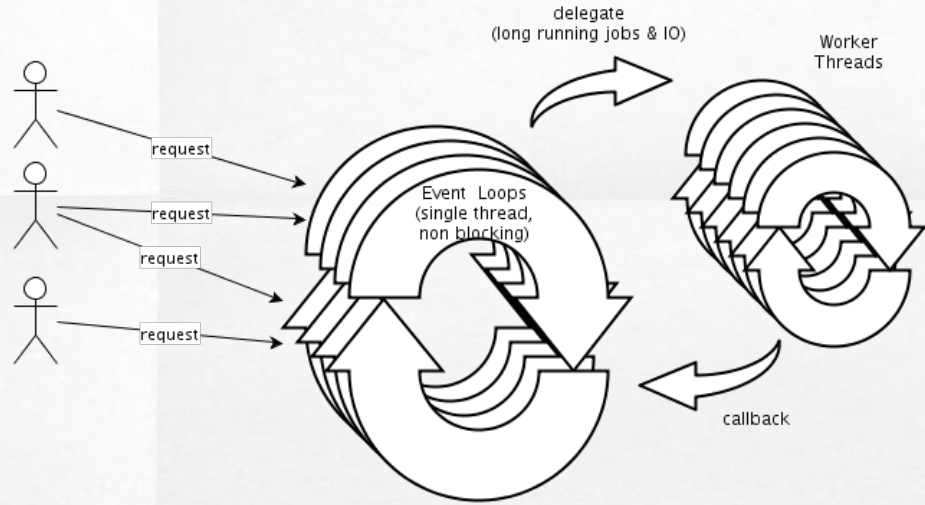
Reactor pattern

- Event Loop
 - non-blocking
 - single thread
 - event handling
- Worker threads
 - concurrency
 - thread pool
 - asynchrony



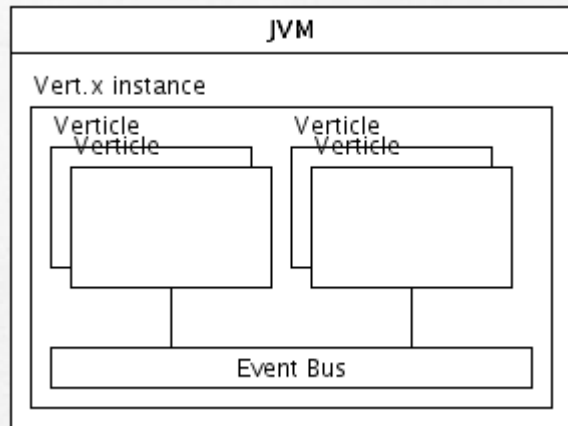
MultiReactor pattern

- Event Loops
 - non-blocking
 - **multiple** threads
 - event handling
- Worker threads
 - concurrency
 - thread pool
 - asynchrony



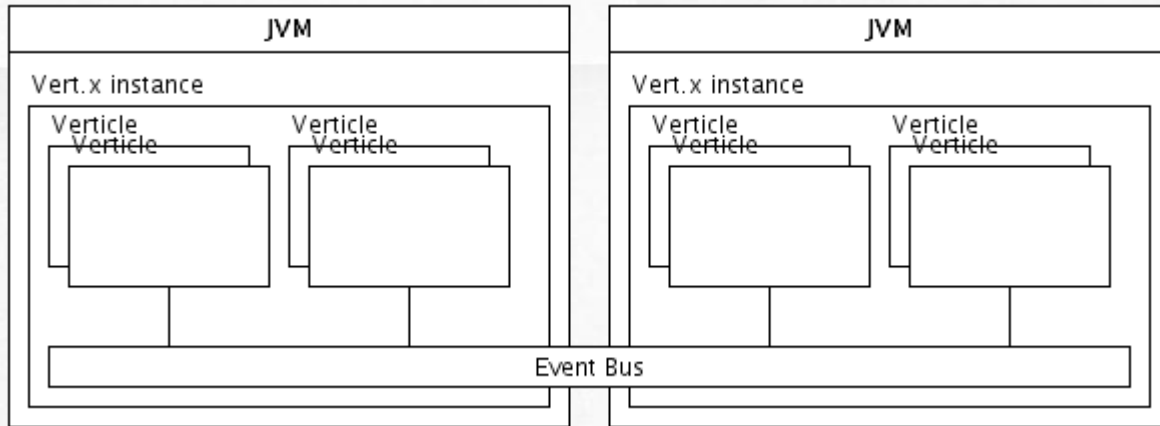
Vert.x overview

- Vert.x instance
- Verticle
- Event Bus

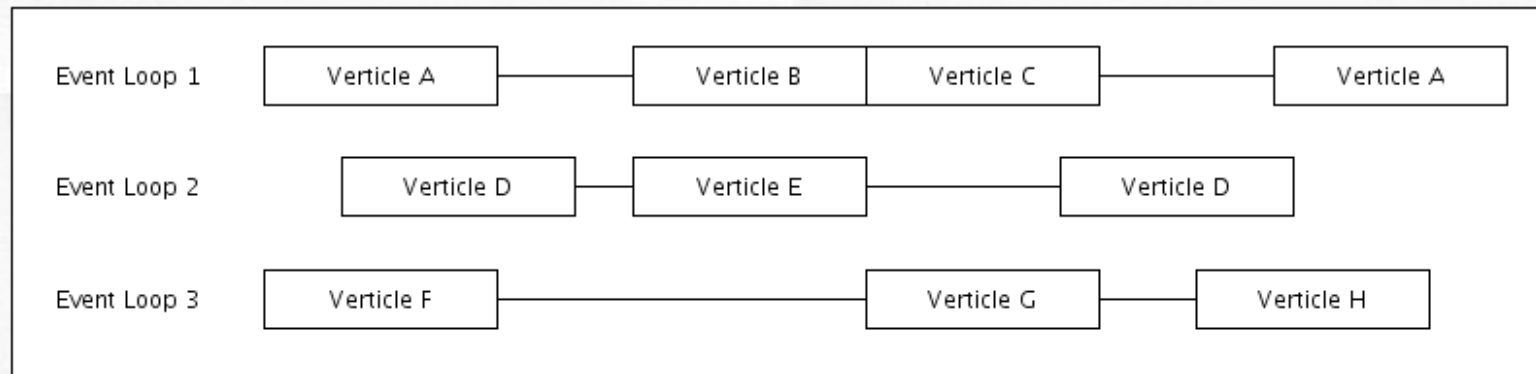


Vert.x clustering

- Event Bus @ Hazelcast



Vert.x loops



Links

- <http://blog.paralleluniverse.co/2014/02/04/littles-law/>
- <http://blog.paralleluniverse.co/2014/05/29/cascading-failures/>
- http://en.wikipedia.org/wiki/Little%27s_law
- http://en.wikipedia.org/wiki/Reactor_pattern